

## UNIT – II

1. \_\_\_\_\_ is a platform for constructing data flows for extract, transform, and load (ETL) processing and analysis of large datasets.

- a) Pig Latin
- b) Oozie
- c) Pig
- d) Hive

2. Point out the correct statement :

- a) Hive is not a relational database, but a query engine that supports the parts of SQL specific to querying data
- b) Hive is a relational database with SQL support
- c) Pig is a relational database with SQL support
- d) All of the mentioned

3. \_\_\_\_\_ hides the limitations of Java behind a powerful and concise Clojure API for Cascading.

- a) Scalding
- b) HCatalog
- c) Cascalog
- d) All of the mentioned

4. Hive also support custom extensions written in :

- a) C#
- b) Java
- c) C
- d) C++

5. Point out the wrong statement :

- a) Elastic MapReduce (EMR) is Facebook's packaged Hadoop offering
- b) Amazon Web Service Elastic MapReduce (EMR) is Amazon's packaged Hadoop offering

- c) Scalding is a Scala API on top of Cascading that removes most Java boilerplate
- d) All of the mentioned

6. \_\_\_\_\_ is the most popular high-level Java API in Hadoop Ecosystem

- a) Scalding
- b) HCatalog
- c) Cascalog
- d) Cascading

7. \_\_\_\_\_ is general-purpose computing model and runtime system for distributed data analytics.

- a) Mapreduce
- b) Drill
- c) Oozie
- d) None of the mentioned

8. The Pig Latin scripting language is not only a higher-level data flow language but also has operators similar to :

- a) SQL
- b) JSON
- c) XML
- d) All of the mentioned

9. \_\_\_\_\_ jobs are optimized for scalability but not latency.

- a) Mapreduce
- b) Drill
- c) Oozie
- d) Hive

10. \_\_\_\_\_ is a framework for performing remote procedure calls and data serialization.

- a) Drill
- b) BigTop
- c) Avro
- d) Chukwa

