

UNIT – IV

1. A _____ node acts as the Slave and is responsible for executing a Task assigned to it by the JobTracker.

- a) MapReduce
- b) Mapper
- c) TaskTracker
- d) JobTracker

2. Point out the correct statement :

- a) MapReduce tries to place the data and the compute as close as possible
- b) Map Task in MapReduce is performed using the Mapper() function
- c) Reduce Task in MapReduce is performed using the Map() function
- d) All of the mentioned

3. _____ part of the MapReduce is responsible for processing one or more chunks of data and producing the output results.

- a) Maptask
- b) Mapper
- c) Task execution
- d) All of the mentioned

4. _____ function is responsible for consolidating the results produced by each of the Map() functions/tasks.

- a) Reduce
- b) Map
- c) Reducer
- d) All of the mentioned

5. Point out the wrong statement :

- a) A MapReduce job usually splits the input data-set into independent chunks which are processed by the map tasks in a completely parallel manner
- b) The MapReduce framework operates exclusively on pairs
- c) Applications typically implement the Mapper and Reducer interfaces to provide the map and reduce methods

d) None of the mentioned

6. Although the Hadoop framework is implemented in Java , MapReduce applications need not be written in :

a) Java

b) C

c) C#

d) None of the mentioned

7. _____ is a utility which allows users to create and run jobs with any executables as the mapper and/or the reducer.

a) Hadoop Strdata

b) Hadoop Streaming

c) Hadoop Stream

d) None of the mentioned

8. _____ maps input key/value pairs to a set of intermediate key/value pairs.

a) Mapper

b) Reducer

c) Both Mapper and Reducer

d) None of the mentioned

9. The number of maps is usually driven by the total size of :

a) inputs

b) outputs

c) tasks

d) None of the mentioned

10. _____ is the default Partitioner for partitioning key space.

a) HashPar

b) Partitioner

c) HashPartitioner

d) None of the mentioned